

.conf2015

Observations and Recommendations on Splunk Performance

Dritan Bitincka
Splunk Technical Services



Disclaimer

During the course of this presentation, we may make forward looking statements regarding future events or the expected performance of the company. We caution you that such statements reflect our current expectations and estimates based on factors currently known to us and that actual events or results could differ materially. For important factors that may cause actual results to differ from those contained in our forward-looking statements, please review our filings with the SEC. The forward-looking statements made in the this presentation are being made as of the time and date of its live presentation. If reviewed after its live presentation, this presentation may not contain current or accurate information. We do not assume any obligation to update any forward looking statements we may make.

In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not, be incorporated into any contract or other commitment. Splunk undertakes no obligation either to develop the features or functionality described or to include any such feature or functionality in a future release.

About me

- Member of Splunk Tech Services
- Large scale deployments
- Cloud and Big Data
- Fifth .Conf

Agenda

- Performance & Bottlenecks
- Understanding Indexing
 - Index-time pipelines
 - Index testing
- Understanding searching in **isolation** & under **indexing load**
 - Types of searches
 - Mixed workload impact on resources

Testing Disclaimers

- Testing on arbitrary datasets in a “closed course” (lab) environment
- Do not take out of context

Typical “*my Splunk is not performing well*” conversation

A: My Splunk is slow.

B: Okay, so what exactly is slow?

A: I dunno, it just feels slow...maybe I'll just get some SSDs.

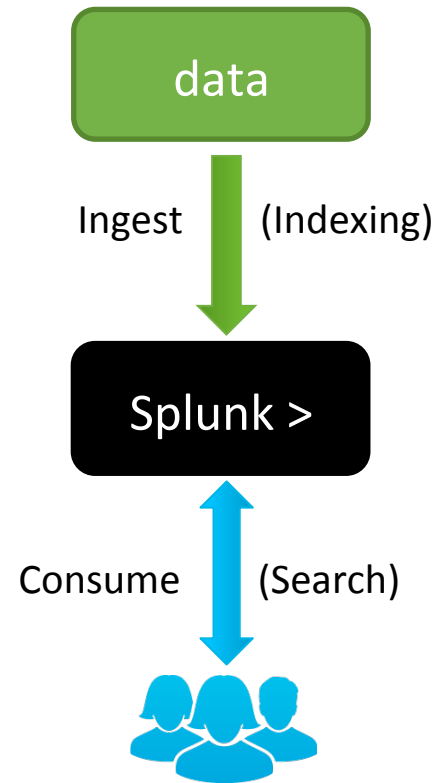


Splunk, like all distributed computing systems, has various bottlenecks that manifest themselves differently depending on workloads being processed.

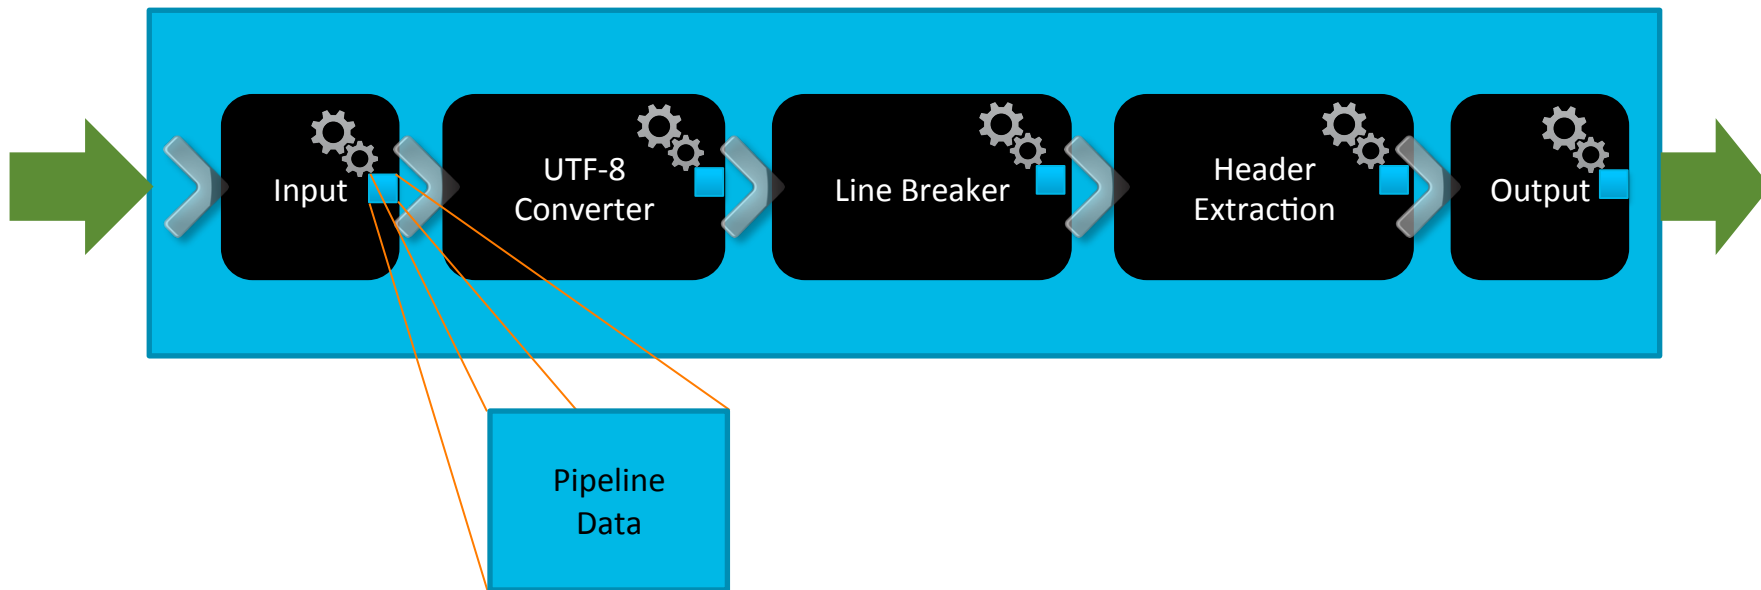
- *Winston Churchill*

Identifying performance bottlenecks

- Understand data flows
 - Splunk operations pipelines
- Instrument
 - Capture metrics for relevant operations
- Run tests
- Draw conclusions
 - Chart and table metrics, looks for emerging patterns
- **Make recommendations**

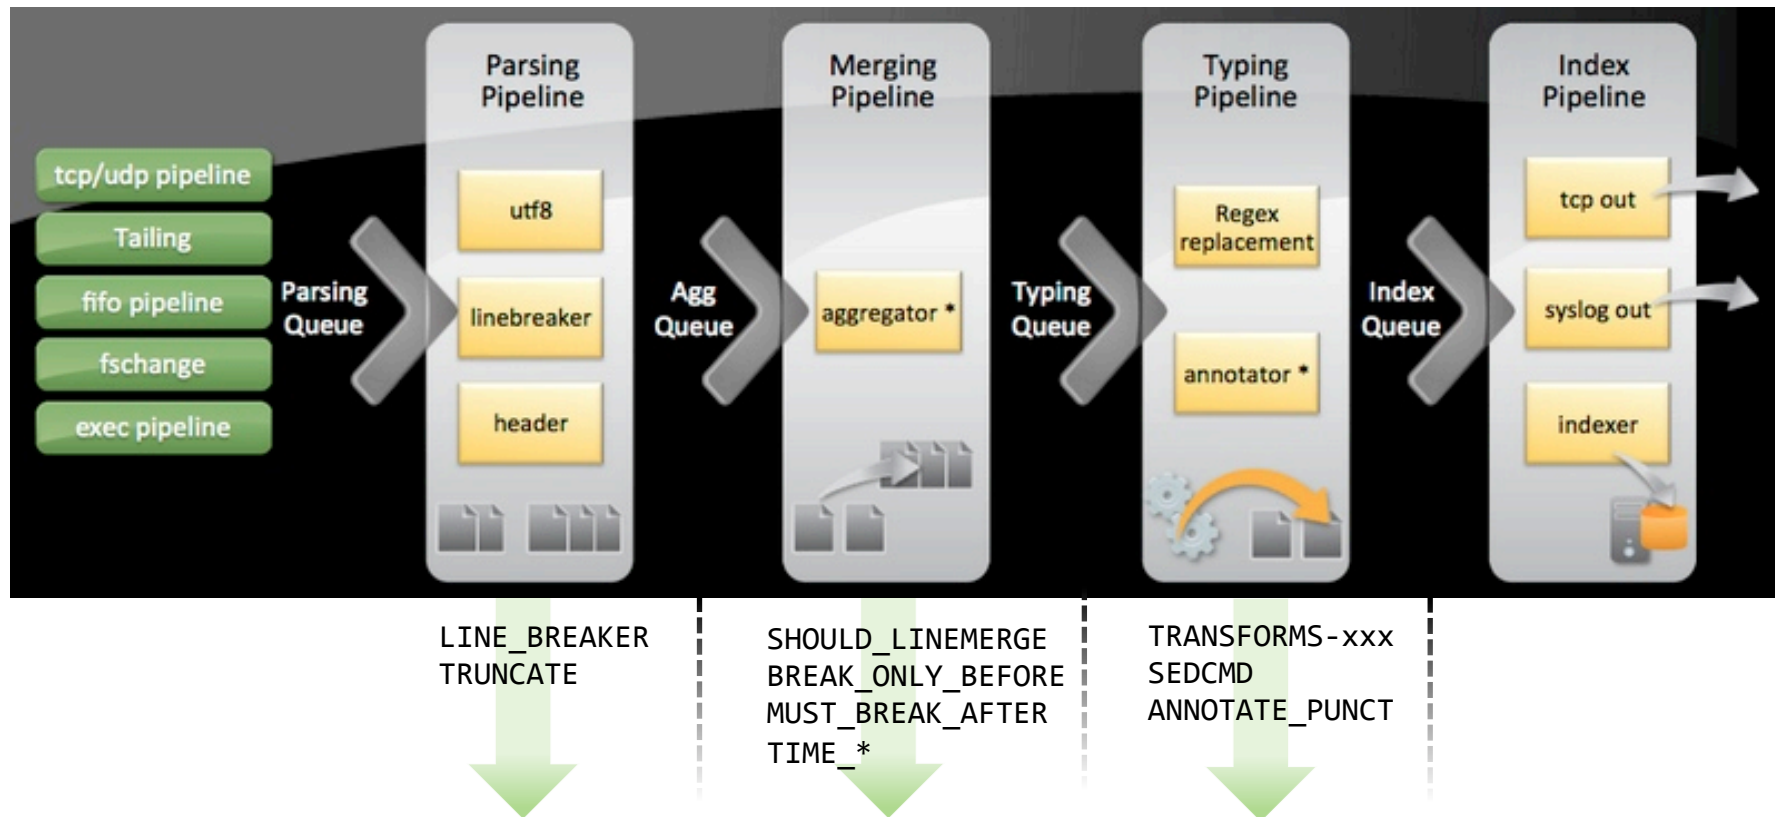


Put that in your pipeline and process it



Splunk data flows thru several such pipelines before it gets indexed

A ton of pipelines



Index-time processing

Event
Breaking

LINE_BREAKER <where to break the stream>

SHOULD_LINEMERGE <enable/disable merging>

Timestamp
Extraction

MAX_TIMESTAMP_LOOKAHEAD <# chars in to look for ts>

TIME_PREFIX <pattern before ts>

TIME_FORMAT <strftime format string to extract ts>

Typing


ANNOTATE_PUNCT <enable/disable punct:: extraction>

Testing: dataset A

- 10M syslog-like events:

```
. . .  
Sat, 06 Apr 2014 15:55:39 PDT <syslog message >  
Sat, 06 Apr 2014 15:55:40 PDT <syslog message >  
Sat, 06 Apr 2014 15:55:41 PDT <syslog message >  
. . .
```


- Push data thru:
 - Parsing > Merging > Typing Pipelines
 - Skip Indexing
 - Tweak various props.conf settings




```
MLA: MAX_TIMESTAMP_LOOKAHEAD = 30  
TP:  TIME_PREFIX = ^  
TF:  TIME_FORMAT = %a, %d %b %Y %H:%M:%S %Z  
LM:  SHOULD_LINEMERGE = false  
AP:  ANNOTATE_PUNCT = false
```

- Measure

Index-time pipeline results






- All pre-indexing pipelines are expensive at default settings.
 - Price of flexibility
- If you're looking for performance, minimize generality
 - `LINE_BREAKER`
 - `SHOULD_LINEMERGE`
 - `MAX_TIMESTAMP_LOOKAHEAD`
 - `TIME_PREFIX`
 - `TIME_FORMAT`

Next: let's index a dataset B


- Generate a much larger dataset (1TB)
 - High cardinality, ~380 Bytes/event, 2.86B events
- Forward to indexer as fast as possible
 - Indexer:
 - 12 CPU@2.67Ghz HT
 - 12GB RAM,
 - 14x15KRPM @146GB/ea
 - No other load on the box
- **Measure**



Indexing: CPU



Indexing: IO



Writes

Reads

Indexing Test Findings

- CPU Utilization
 - ~**35%** in this case, **4-5** Real CPU Cores
- IO Utilization
 - Characterized by both reads and writes but not as demanding as search.
Note the *splunk-optimize* process.
- Ingestion Rate
 - **22MB/s**
 - “Speed of Light” – no search load present on the server

Index Pipeline Parallelization

- Splunk 6.3+ can maintain multiple independent pipelines sets
 - i.e. same as if each set was running on its own indexer
- If machine is under-utilized (CPU and I/O), you can configure the indexer to run **2** such sets.
- Achieve roughly **double** the indexing throughput capacity.
- Try not to set over **2**
- Be mindful of associated resource consumption

Indexing Test Conclusions

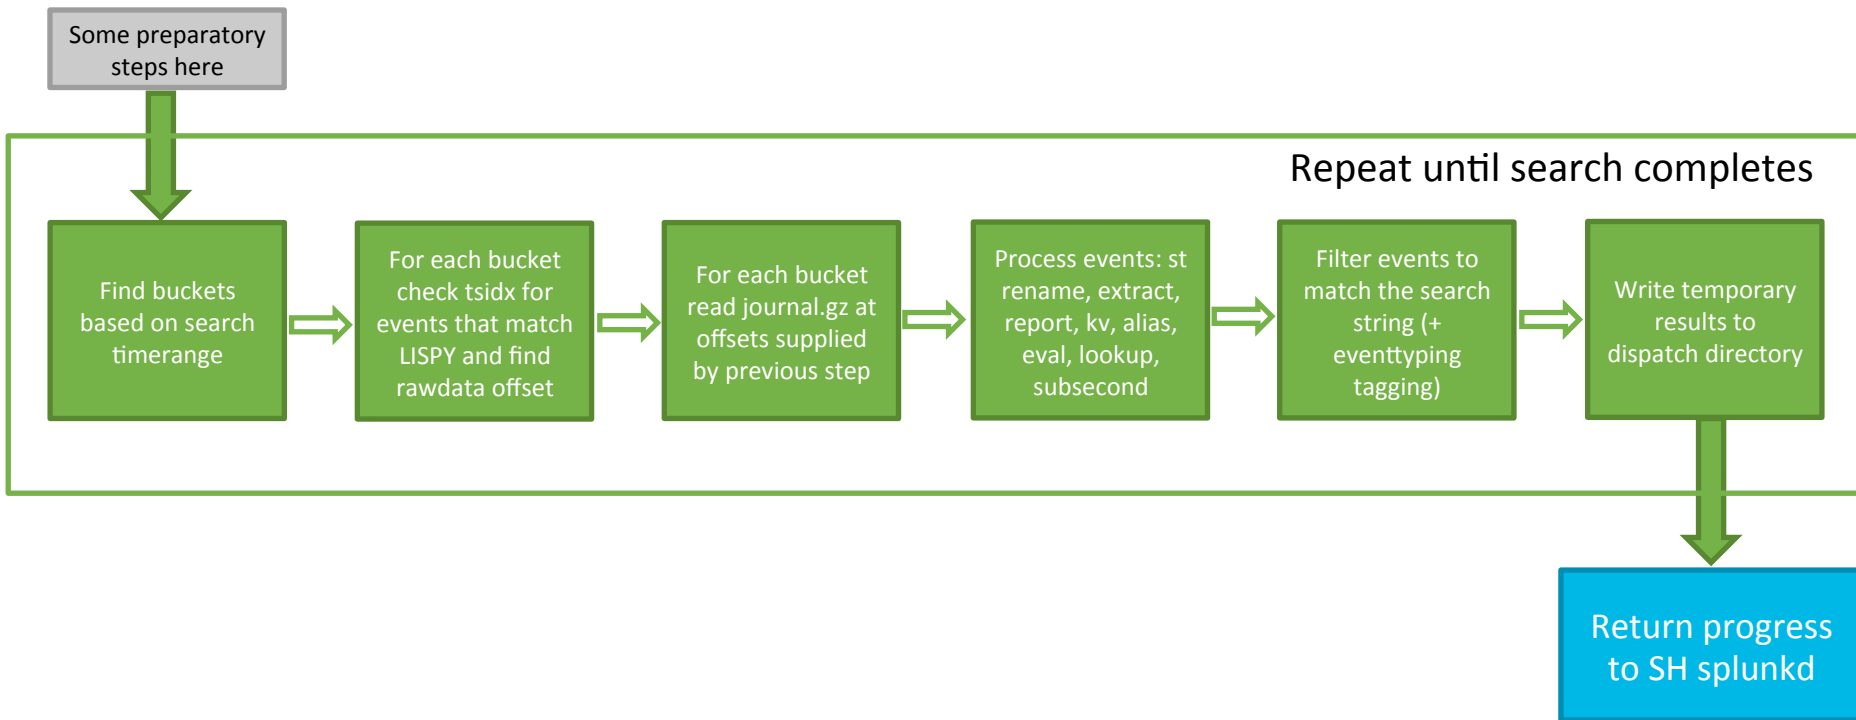
- **Distribute** as much as you can – Splunk scales horizontally
- Enable more pipelines but be aware of compute tradeoff
- **Tune** event **breaking** and **timestamping** attributes in props.conf whenever possible
- Faster disk (ex. SSDs) would not have necessarily improved indexing throughput by much
- Faster, but not more, CPUs would have improved indexing throughput – modulo pipeline parallelization

Next: Searching

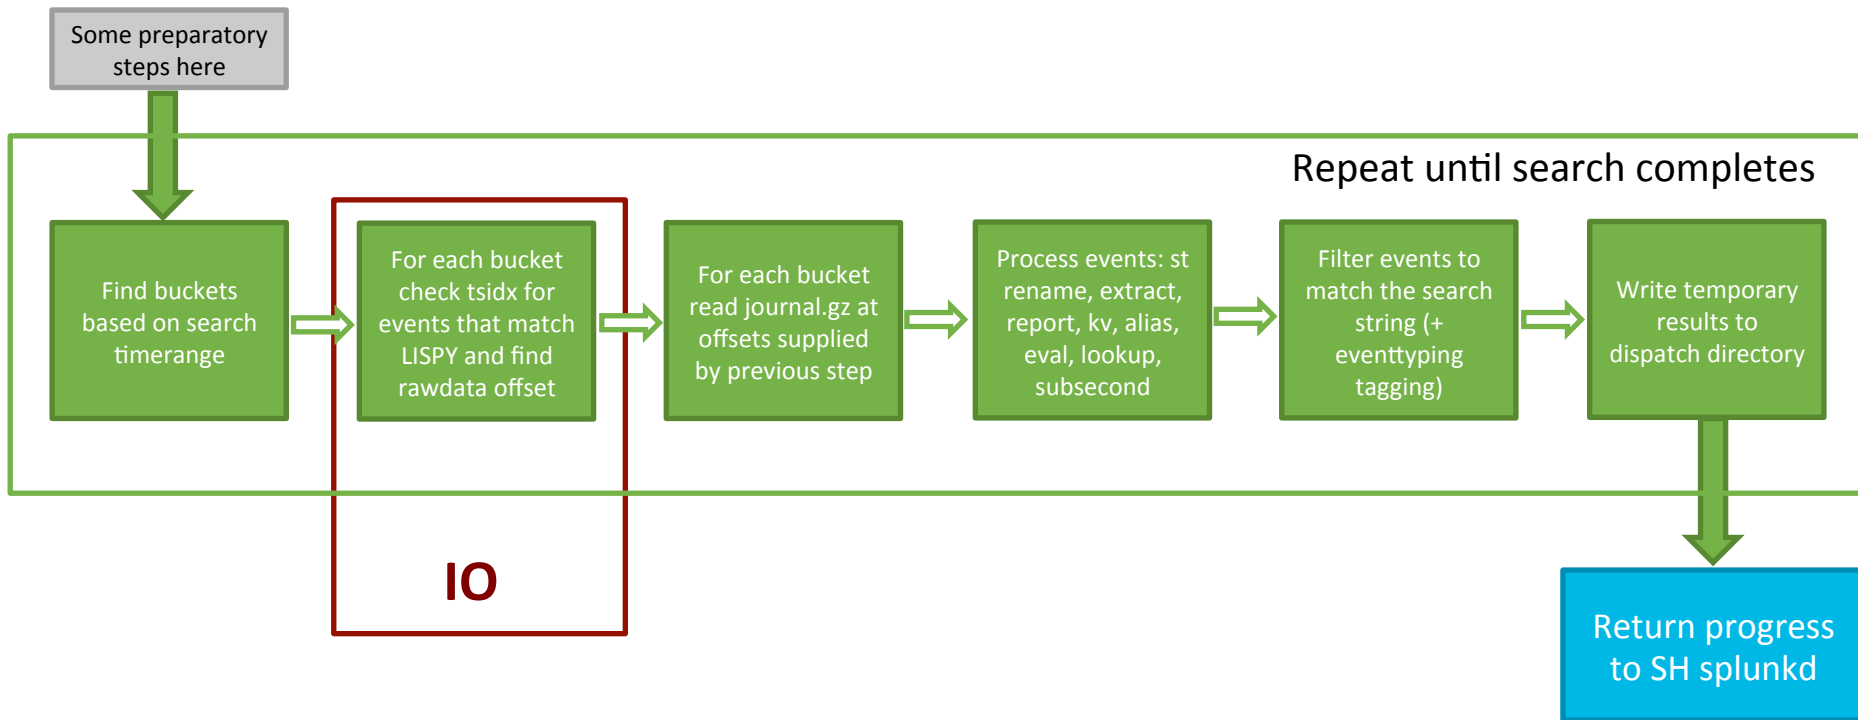
- Real-life search workloads are extremely complex and very varied to be profiled correctly
- But, we can generate arbitrary workloads covering a wide spectrum of resource utilization and profile those instead. Actual profile will fall somewhere in between.



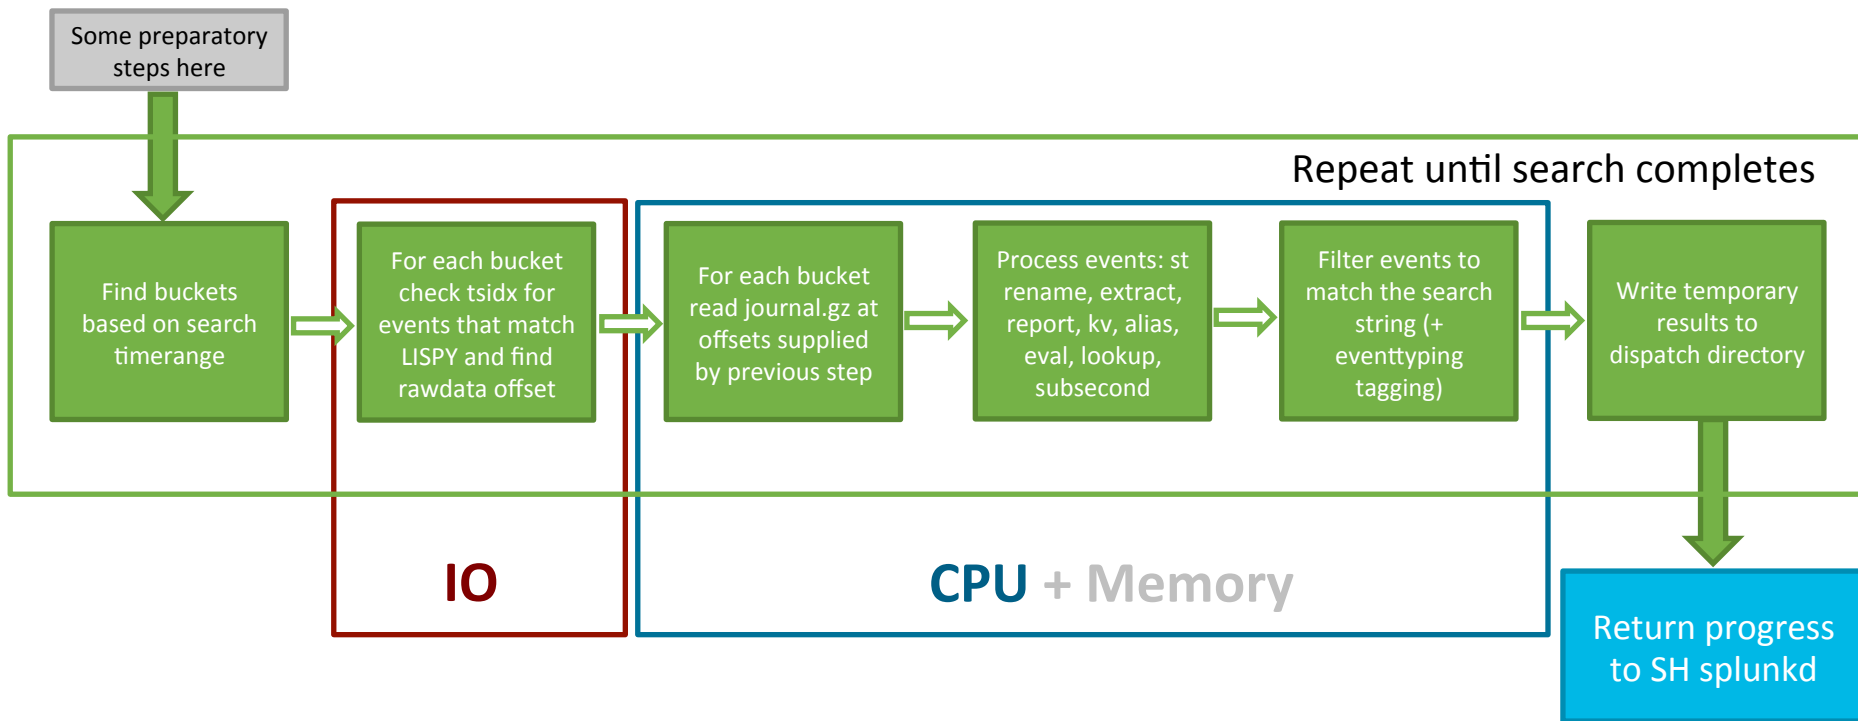
Search pipeline (High Level)



Search pipeline boundedness



Search pipeline boundedness



Search Types


- **Dense**
 - Characterized predominantly by returning **many events** per bucket
`index=web | stats count by clientip`
- **Sparse**
 - Characterized predominantly by returning **some events per bucket**
`index=web some_term | stats count by clientip`
- **Rare**
 - Characterized predominantly by returning **only a few** events per index
`index=web url=onedomain* | stats count by clientip`

Okay, let's test some searches


- Use our already indexed data
 - It contains **many** unique terms with predictable term density
- Search under several term densities and concurrencies
 - Term density: 1/100, 1/1M, 1/100M
 - Search Concurrency: 4 – 60
 - Searches:
 - **Rare: over all 1TB dataset**
 - **Dense: over a preselected time range**
- Repeat all of the above while under an indexing workload
- **Measure**

Dense Searches

% CPU Util. vs. Concurrency (1/100)




% IO Wait vs. Concurrency (1/100)




What's going on here?

Dense Searches

% CPU Util. vs. Concurrency (1/100)



% IO Wait vs. Concurrency (1/100)




What's going on here?


CPU
Bound

Dense Searches with Indexing

% CPU Util. vs. Concurrency (searching 1/100)



Indexing Thruput vs. Concurrency (searching 1/100)



Hitting 100%
earlier

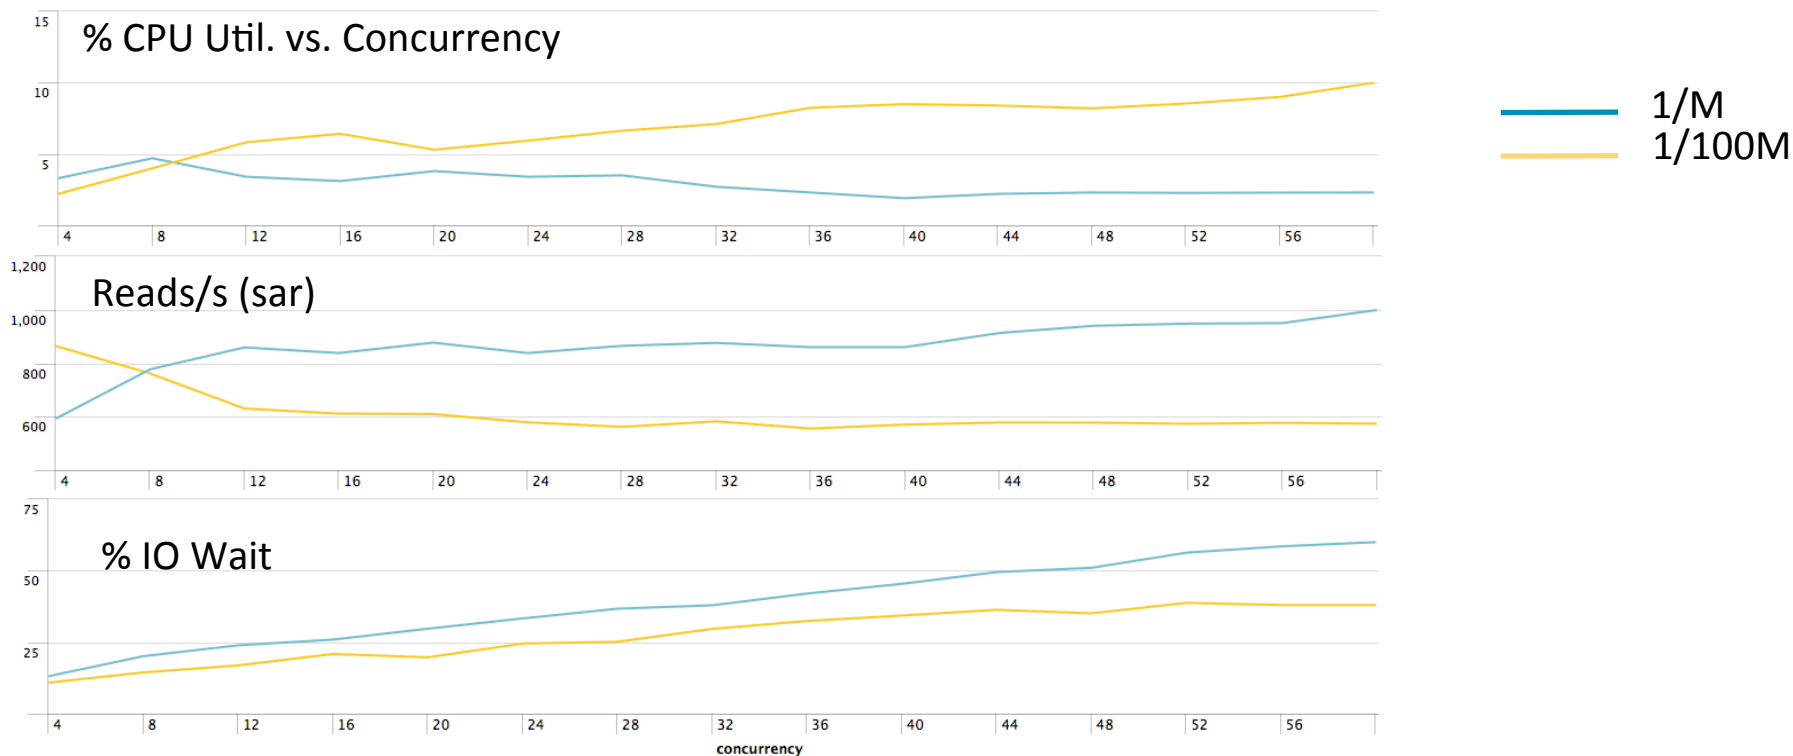
SpeedOfLight
1/100

Test Type	Indexing off	Indexing on	diff
Avg. Search Time	40.58s	43.85s	+8.6%

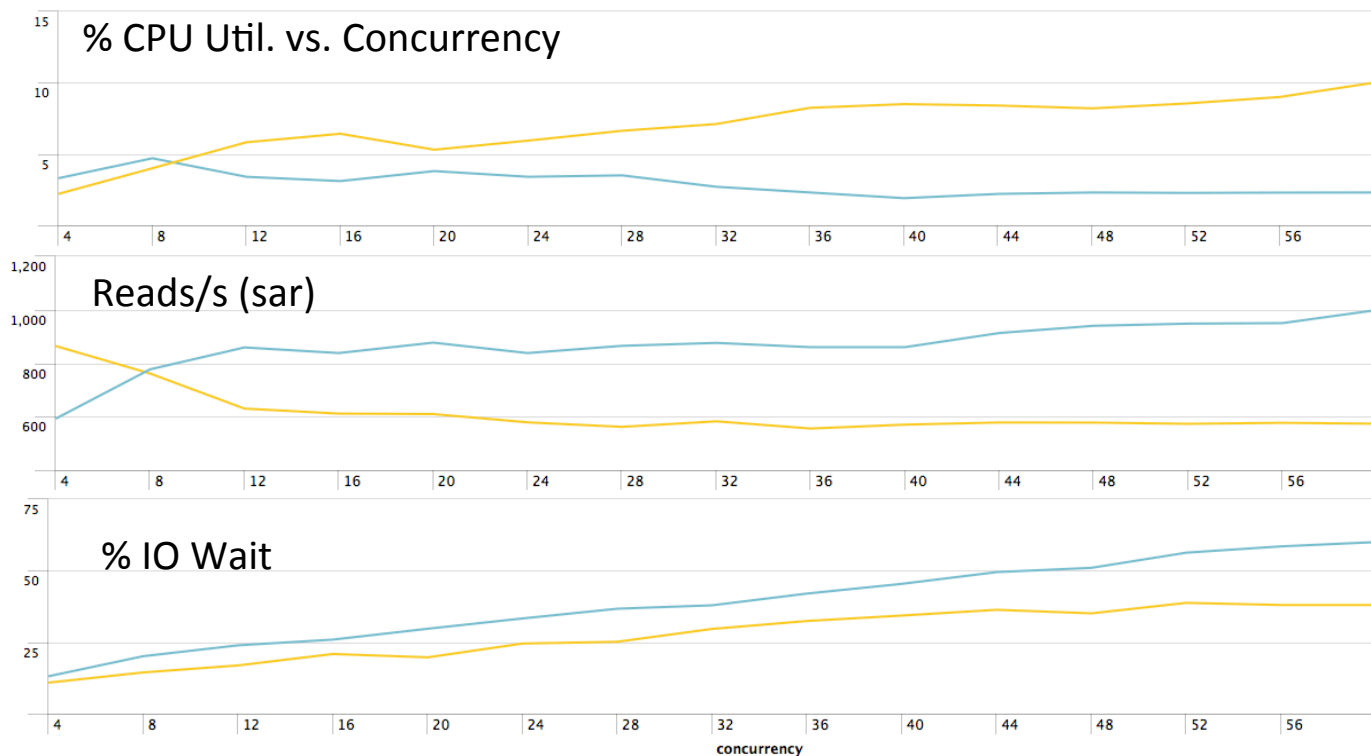
Dense Search Test Conclusions

- Dense workloads are CPU bound
- Dense workloads (reporting, trending etc.) play relatively “okay” with indexing
 - While indexing throughput decreases by ~40% search time increases only marginally
- **Faster disk wont necessarily help as much here**
 - Majority of time in dense searches is spent in CPU decompressing rawdata + other SPL processing
- **Faster and more CPUs would have improved overall performance**


Rare Searches



Rare Searches




Rare Searches with Indexing



More numbers

Indexing Thruput (KB/s) vs. Concurrency

SpeedOfLight
1/1M
1/100M

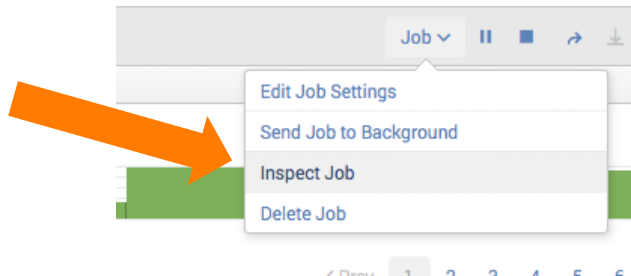


Density	1/1M			1/100M		
Test Type	Indexing off	Indexing on	diff	Indexing off	Indexing on	diff
Avg Search Time	1041s	1806s	+73.5%	231s	304s	+31.6%

Rare Search Conclusions

- Rare workloads (investigative, ad-hoc) are IO bound
- Rare workloads (high IO wait) do not play well with indexing
 - Search time increases significantly when indexing is on. Also, indexing throughput takes an equal hit on the opposite direction.
- 1/100M searches have a lesser impact on IO than 1/1M.
- When indexing is on, in 1/1M case search time increases more substantially vs. 1/100M. Search and indexing are both contenting for IO.
- In case of 1/100M, **bloomfilters** save precious IO which, in turn, allows for a better indexing throughput.
 - **Bloomfilters** are special data structures that indicate with 100% certainty that a term **does not exist** in a bucket (effectively telling the search process to skip that bucket).
 - Note the higher CPU consumption during 1/100M searches with indexing on
- **Faster disks would have definitely helped here**
- **More CPUs would not have improved performance by much**

Is my search CPU or IO bound?



Guideline in absence of full instrumentation

- **command.search.rawdata** ~ CPU Bound
 - Others: .kv, .typer, .calcfields,
- **command.search.index** ~ IO Bound

Search job inspector

This search has completed and has returned 1 result by scanning 4,159,473 events in 20.706 seconds.

The following messages were returned by the search subsystem:

```
DEBUG: Disabling timeline and fields picker for reporting search due to adhoc_search_level=smart
DEBUG: base lisp: [ AND index::internal ]
DEBUG: search context: user="admin", app="aws_app", bs-pathname="/opt/splunk61/etc"
```

(SID: 1410010633.156)

Execution costs

Duration (seconds)		Component	Invocations	Input count	Output count
	0.344	command.addinfo	344	4,159,473	4,159,473
	0.343	command.fields	344	4,159,473	4,159,473
████	7.133	command.prestats	344	4,159,473	343
████████	13.247	command.search	344	-	4,159,473
██████████	10.254	★ command.search.rawdata	344	-	-
	0.363	command.search.kv	343	-	-
	0.344	command.search.tags	344	4,159,473	4,159,473
	0.344	command.search.typer	344	4,159,473	4,159,473
	0.343	command.search.calcfields	343	4,159,473	4,159,473
	0.343	command.search.fieldalias	343	4,159,473	4,159,473
	0.343	command.search.lookups	343	4,159,473	4,159,473
	0.11	command.search.summary	344	-	-
	0	command.search.index.usec_1_8	22	-	-
	0	command.search.index.usec_512_4096	84	-	-
	0	command.search.index.usec_64_512	314	-	-
	0	command.search.index.usec_8_64	116	-	-
	0.345	command.stats.execute_input	345	-	-

Top Takeaways/Re-Cap

- **Indexing**


- **Distribute** – Splunk scales horizontally
- **Tune** event breaking and timestamp extraction
- **Faster** CPUs will help with indexing performance

- **Searching**

- **Distribute** – **Splunk scales horizontally**
- **Dense Search Workloads**
 - CPU Bound, better with indexing than rare workloads
 - Faster and more CPUs will help
- **Rare Search Workloads**
 - IO Bound, not that great with indexing
 - Bloomfilters help significantly
 - Faster disks will help

- **Performance**

- Avoid generality, optimize for expected case and add hardware whenever you can



Use case	What Helps?
Trending, reporting over long term etc.	More distribution Faster, more CPUs
Ad-hoc analysis, investigative type	More distribution Faster Disks, SSDs

Testing Disclaimer Reminder

1. Testing conducted on arbitrary datasets
2. “closed course” (lab) environment
3. Not to be interpreted out of context

.conf2015

You may also like

Architecting and Sizing Your Splunk Deployment

Onboarding Data Into Splunk

Splunk Health Check. How is Your Environment Feeling?

Q & A

THANK YOU

Feedback: dritan@splunk.com

splunk>